




EVAL-DRIVEN DEVELOPMENT (RAGAS)

BUILDING AN AI APPLICATION IS EASY. BUT HOW DO YOU KNOW IF
ITS ANSWERS ARE ACTUALLY GOOD?





WHAT IS AI EVALUATION?

- Measures AI response quality
- Checks if answers are accurate
- Detects hallucinations
- Compares AI performance over time
- Essential for production AI systems





WHAT IS EVAL-DRIVEN DEVELOPMENT?

- Build → Test → Evaluate → Improve
- Measure AI before deploying
- Data-driven AI improvement
- Continuous quality monitoring
- Industry best practice for LLM apps





WHAT IS RAGAS?

- Open-source RAG evaluation framework
- Evaluates retrieval and generation quality
- Measures multiple AI metrics
- Helps improve RAG systems
- Widely used in production AI projects





EVALUATION PIPELINES

- Automated AI quality testing
- Run multiple test questions
- Compare expected vs actual answers
- Generate evaluation reports
- Track improvements over time





HALLUCINATION SCORING

- Detects fabricated information
- Checks if answers match retrieved context
- Reduces misinformation
- Builds trustworthy AI systems
- Critical for enterprise applications





RESPONSE QUALITY METRICS

- Faithfulness
- Answer Relevancy
- Context Precision
- Context Recall
- Response Correctness
- Overall Quality Score





WHY RAGAS MATTERS?

- Improve AI reliability
- Reduce hallucinations
- Better user experience
- Production-ready AI quality
- Enterprise standard evaluation

